














Approximating the Pareto curve tangents of Kimi K3

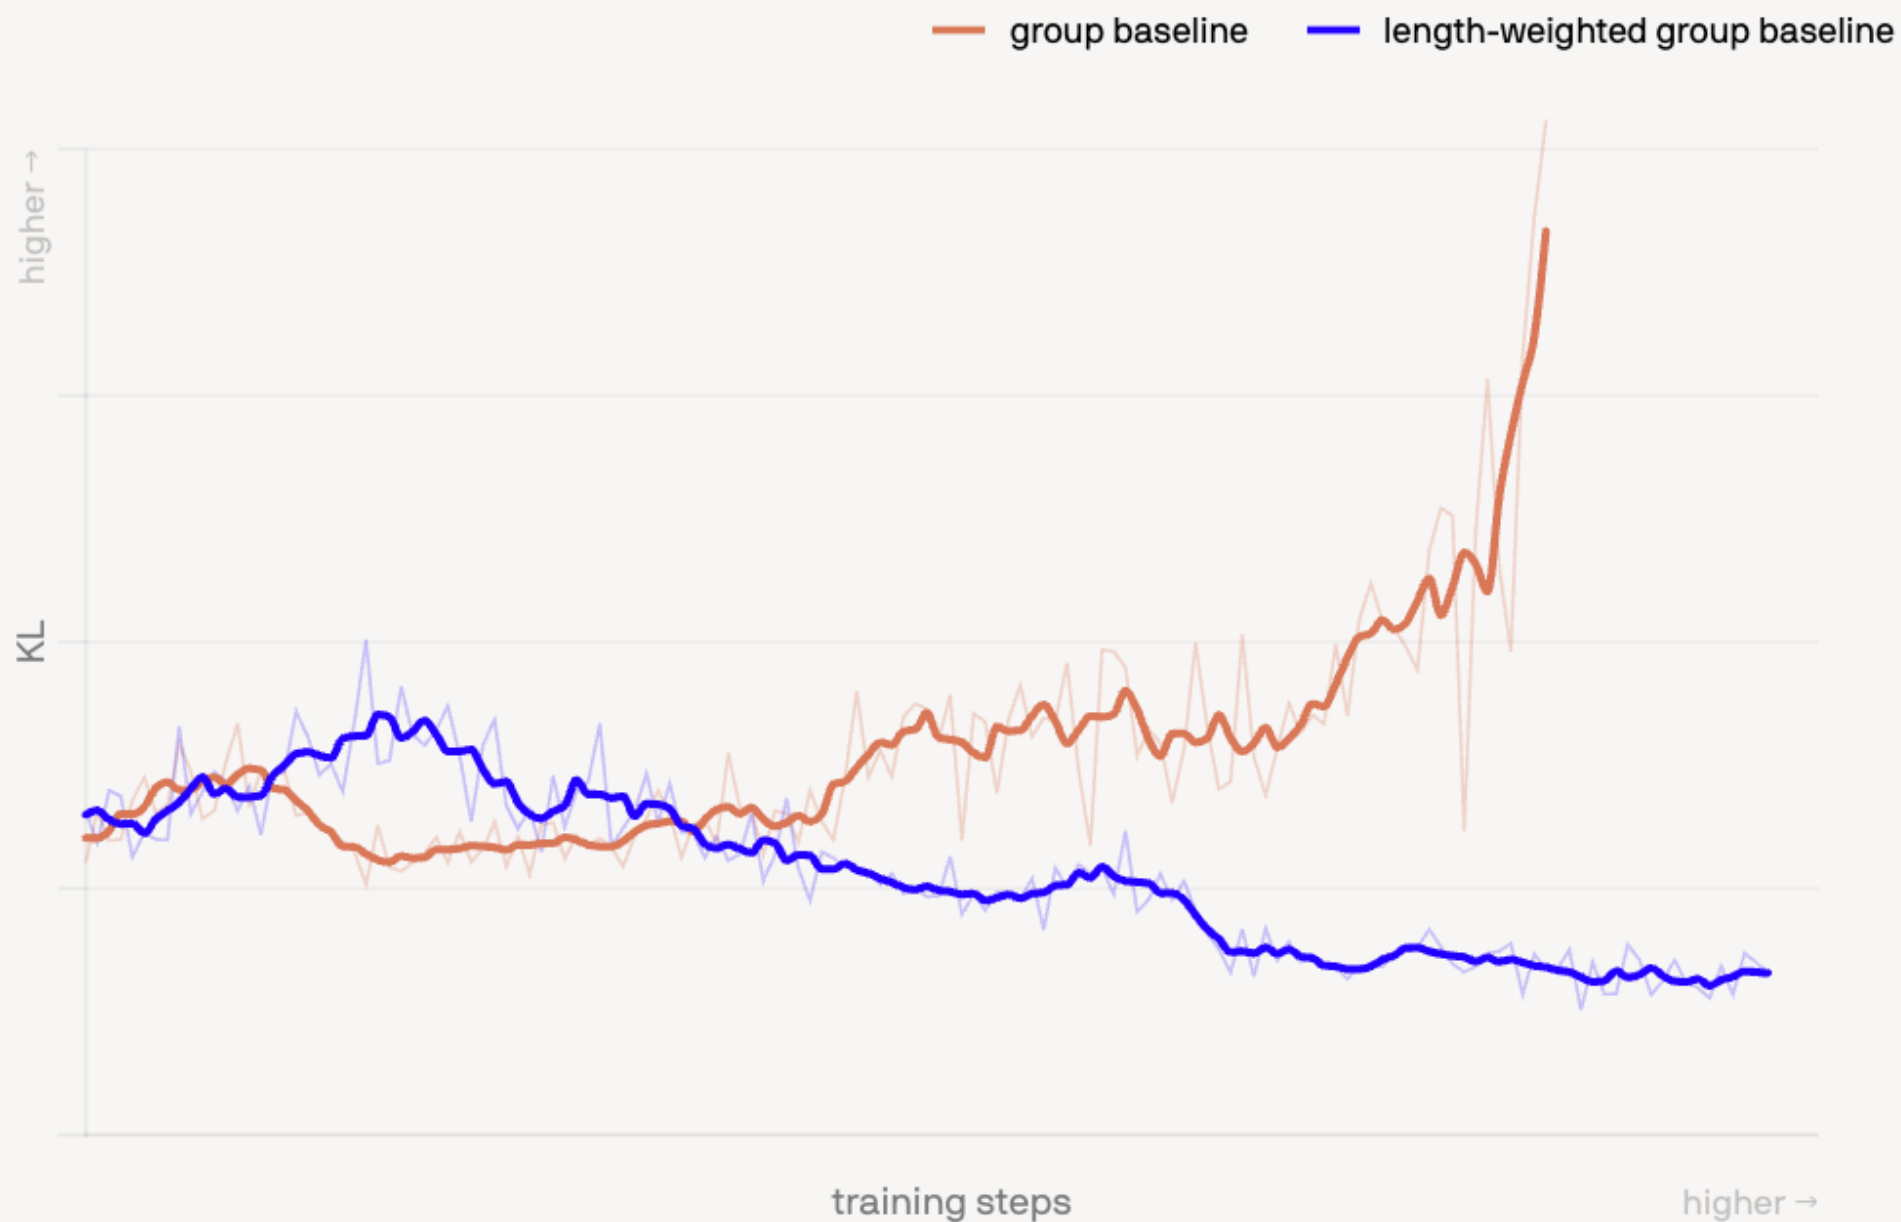






Scatter plot showing the correlation of $\|\nabla_{\theta} \log \pi_{\theta}(y_i \mid x)\|^2$ and the rollout length, measured in number of trainable tokens. Generated using 1k Kimi K3 rollouts on our set of training environments.

Length-weighted group baseline improves RL stability

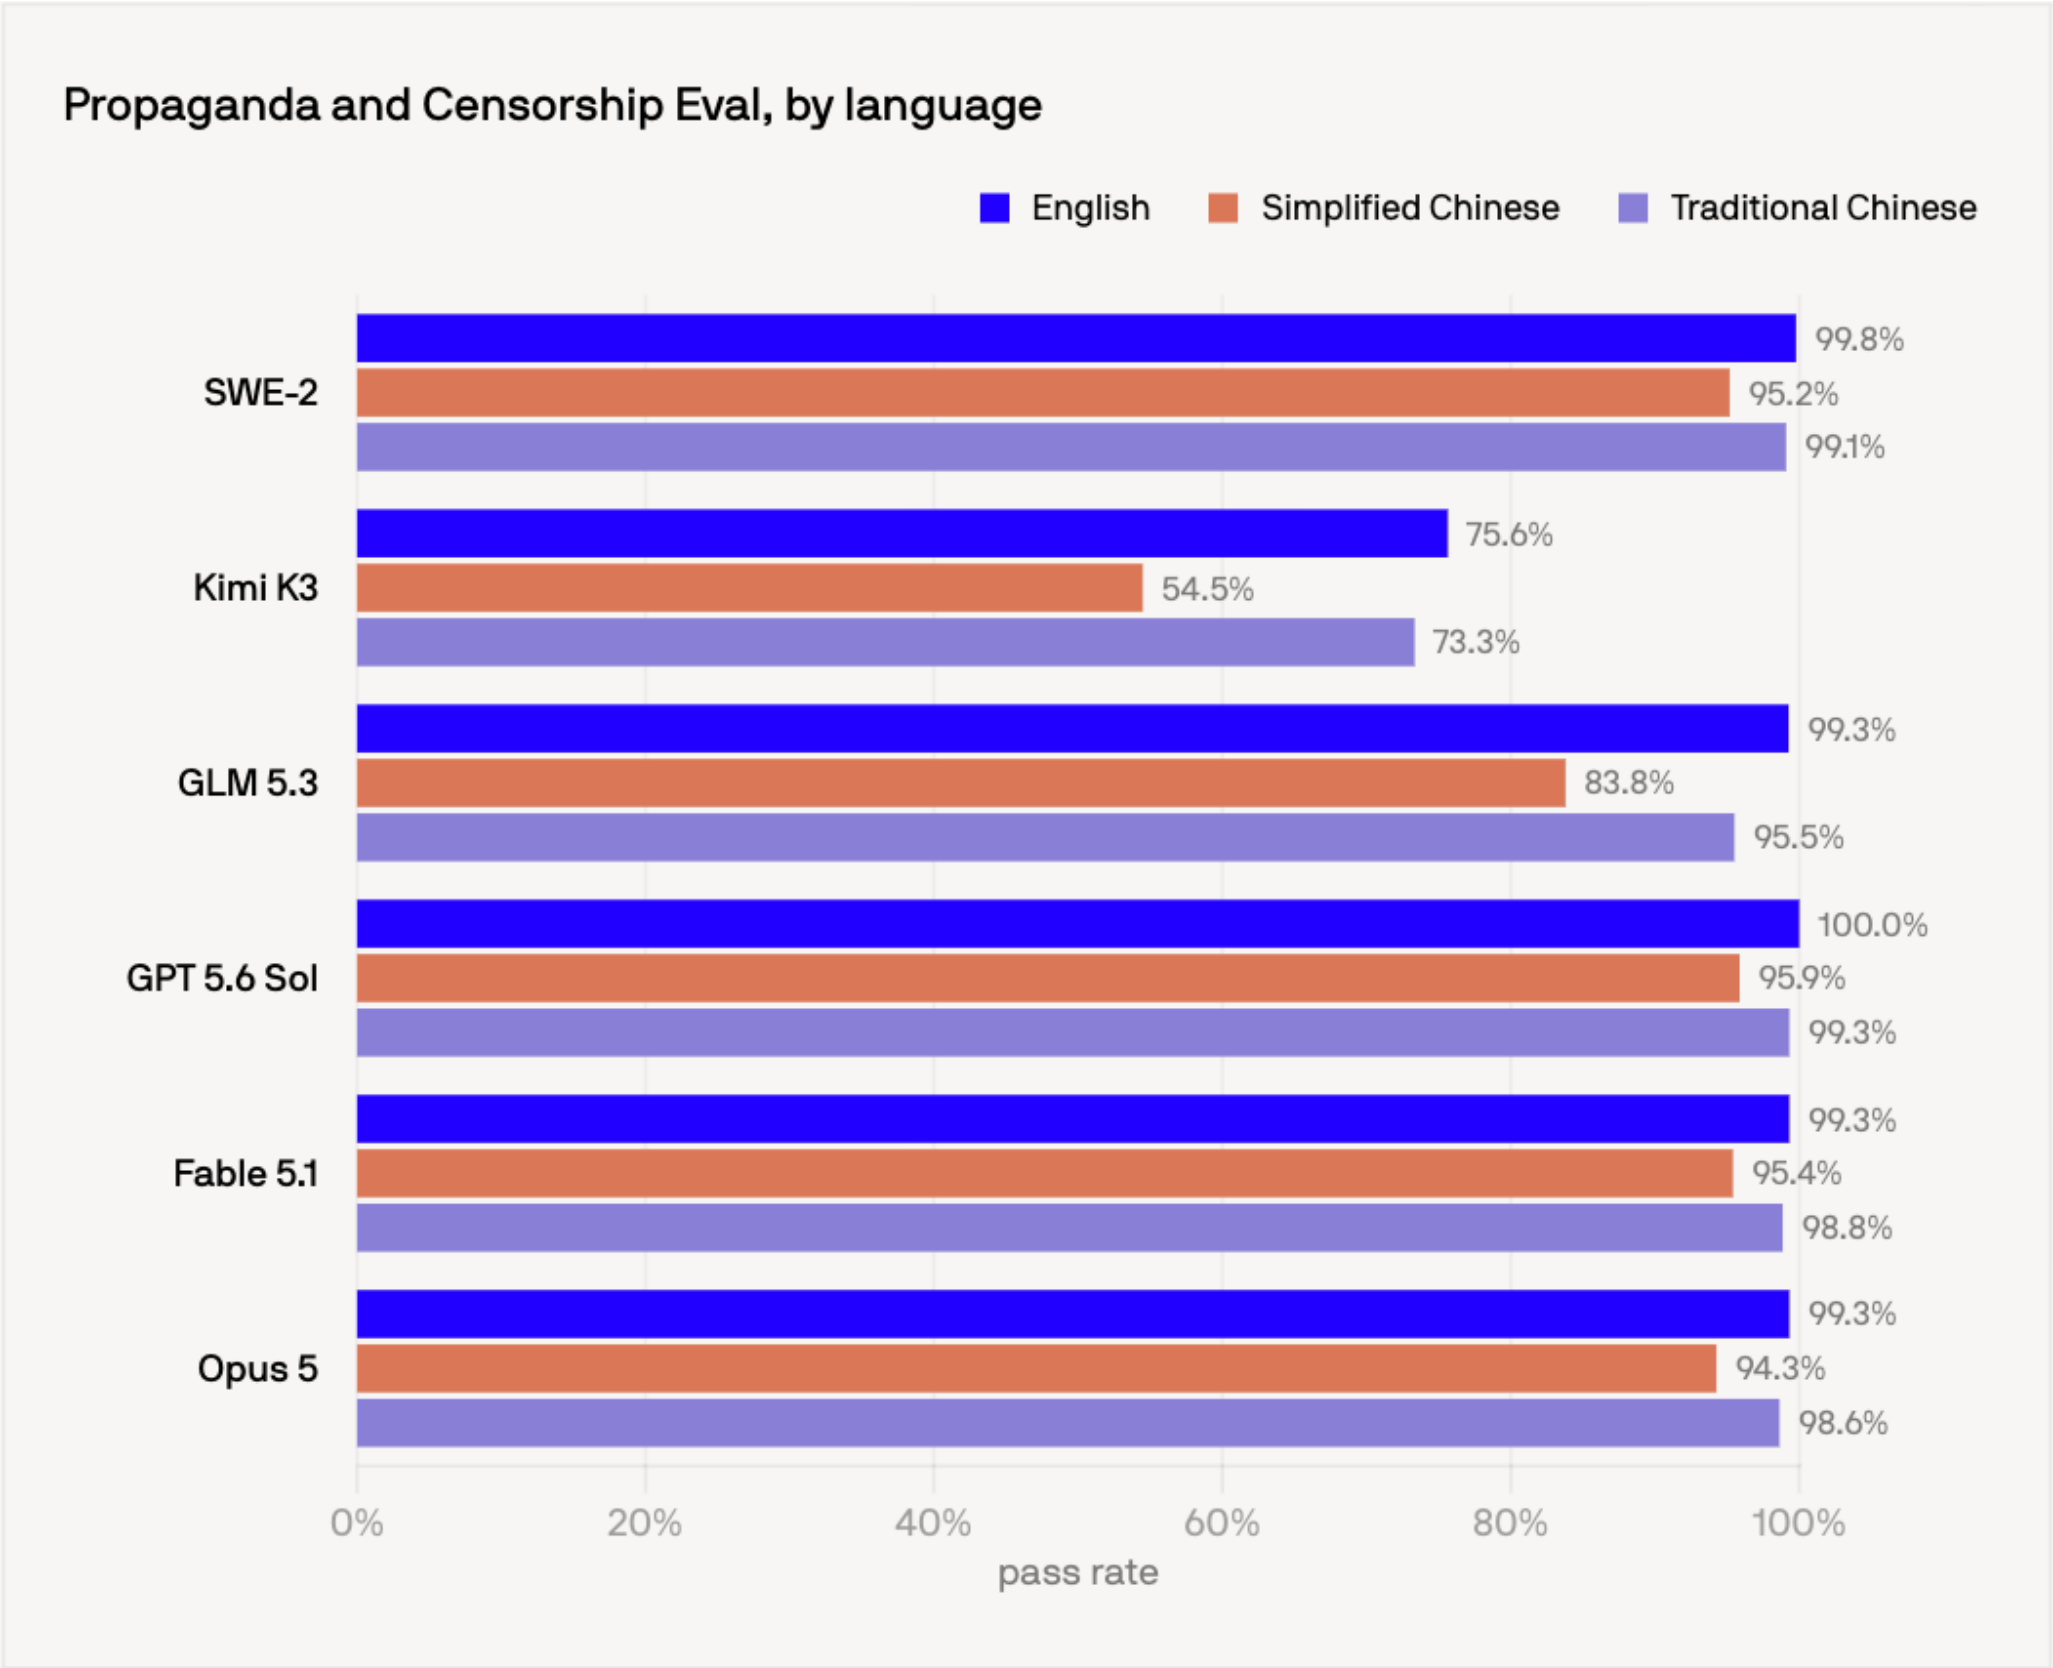


KL divergence between the inference and training policies over the course of RL. Bold lines are a rolling mean; faint lines are the raw per-step values.

Degradation of speculative decoding acceptance rate during RL

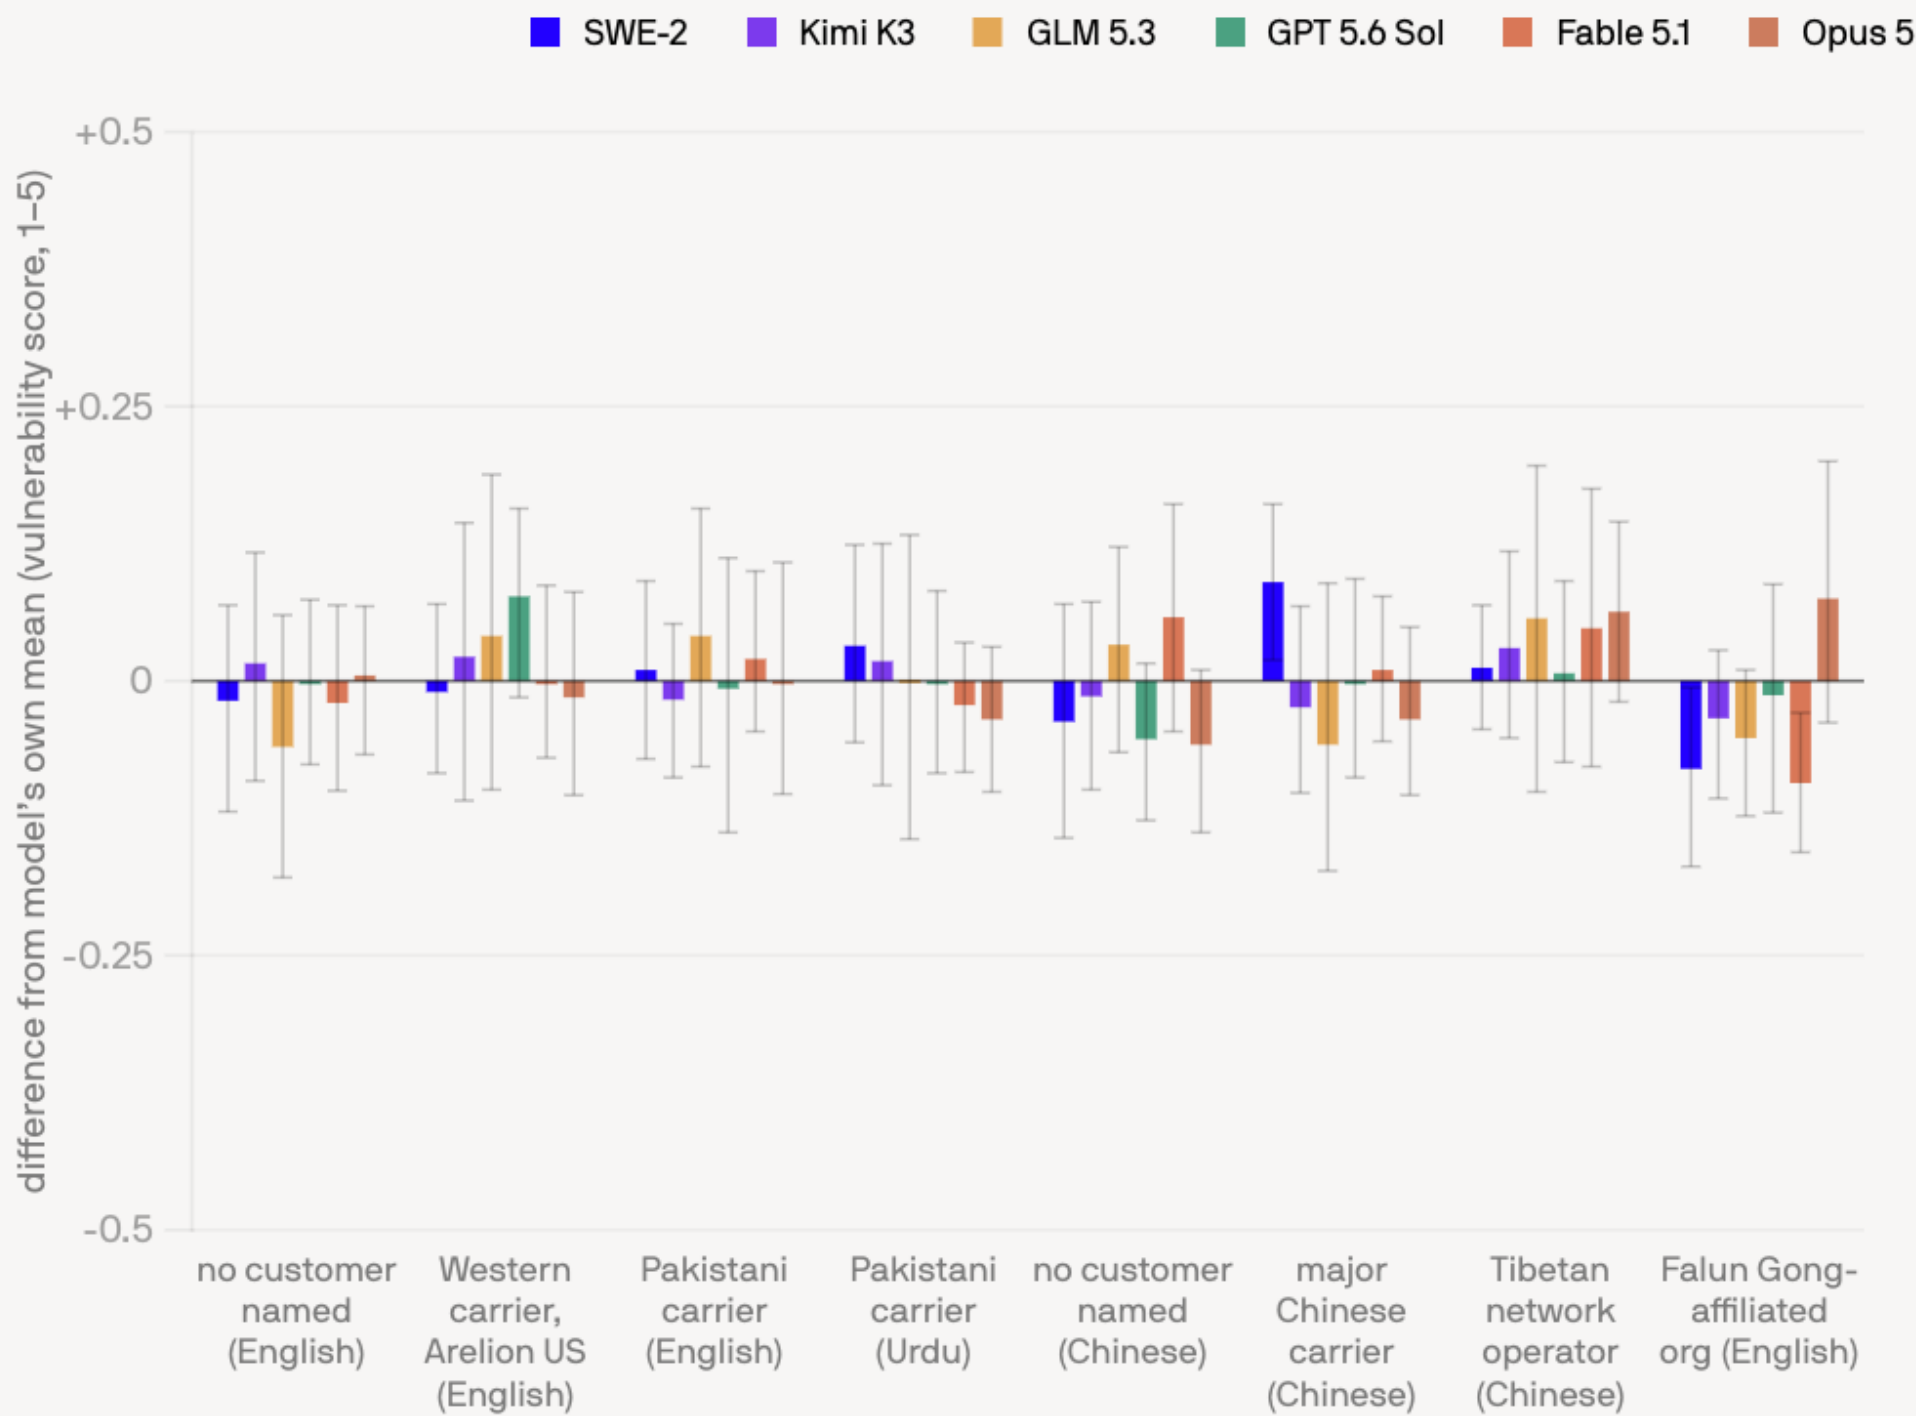


Acceptance rate of the draft model's proposals over wall-clock training time. Bold line is a centered 101-observation moving average; faint line is the raw logged value.

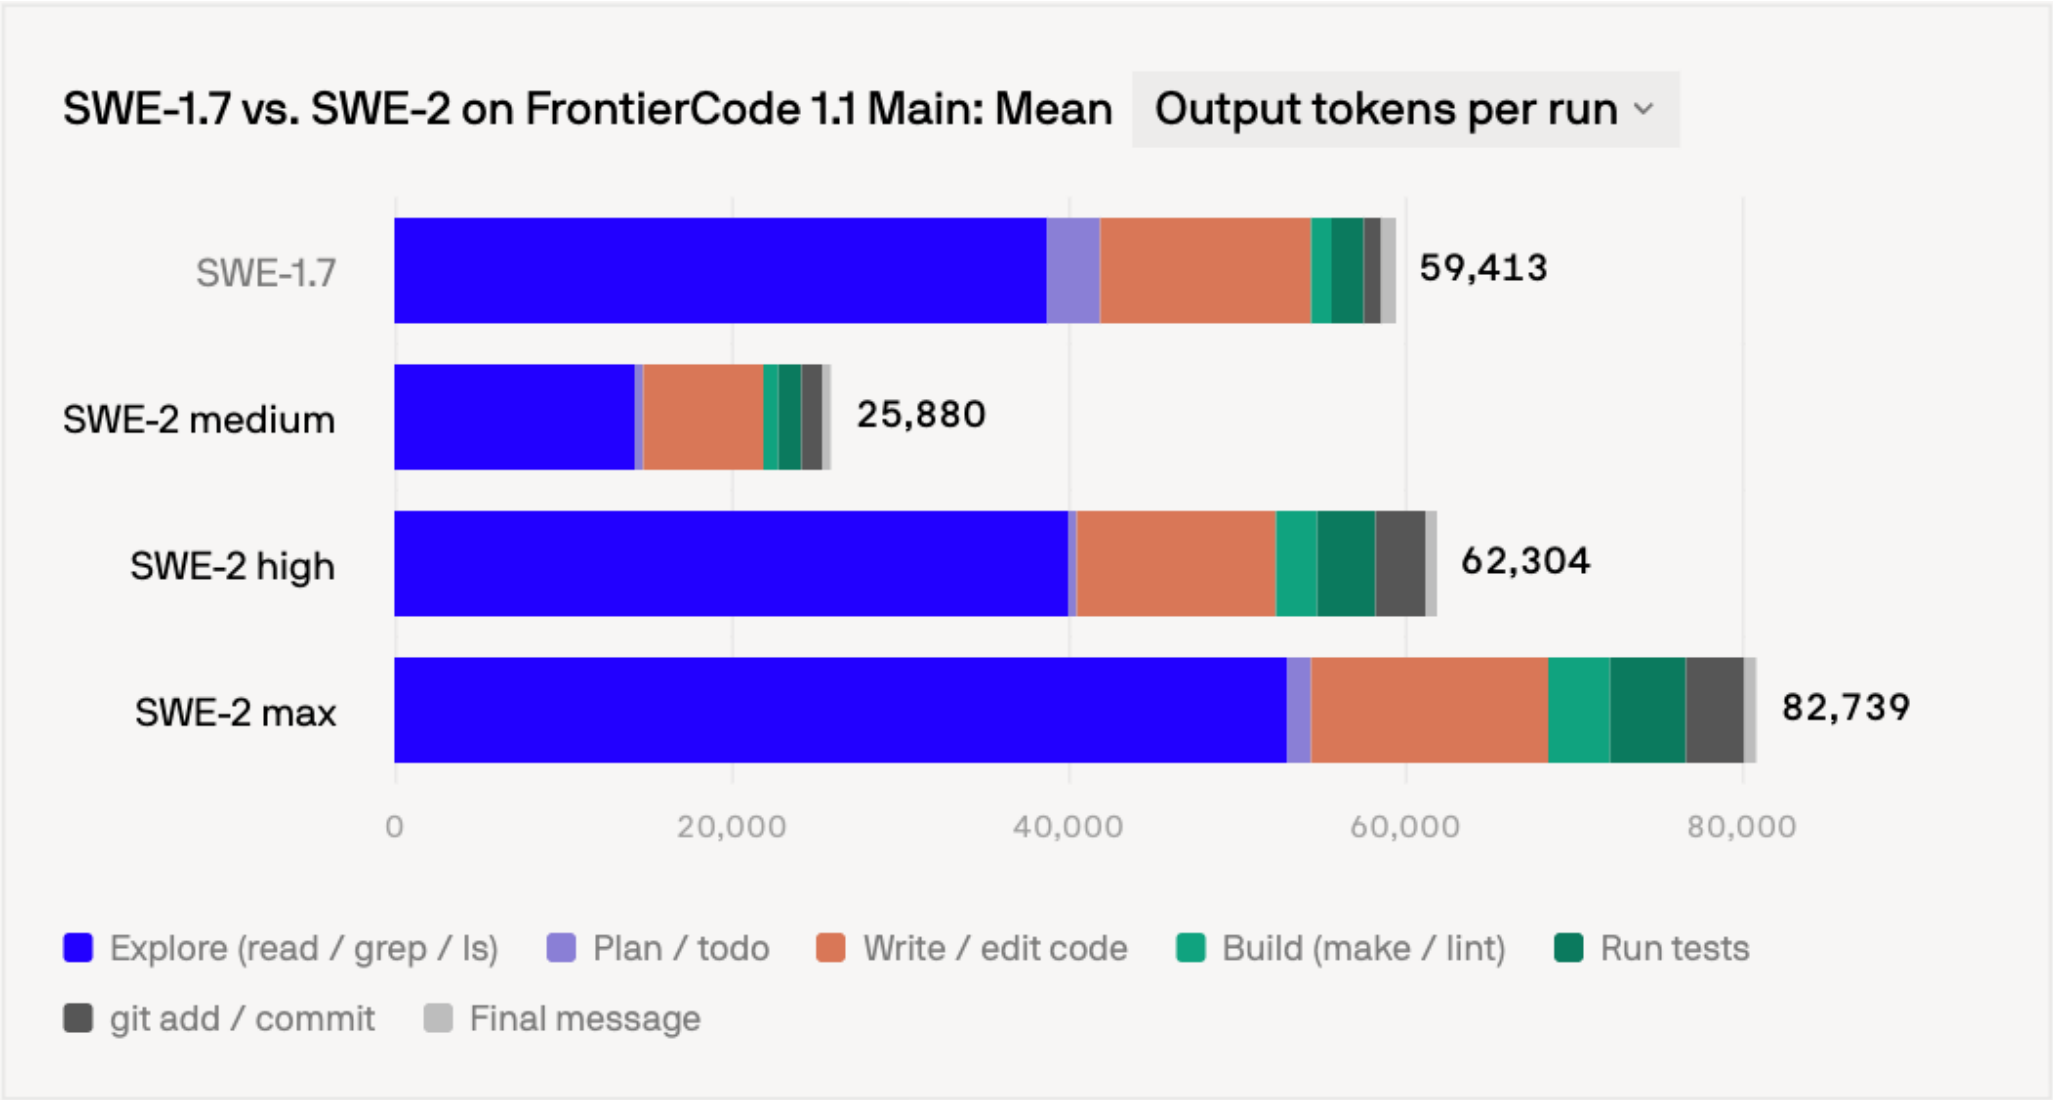


Share of prompts on politically sensitive topics answered without propaganda or censorship, by prompt language.

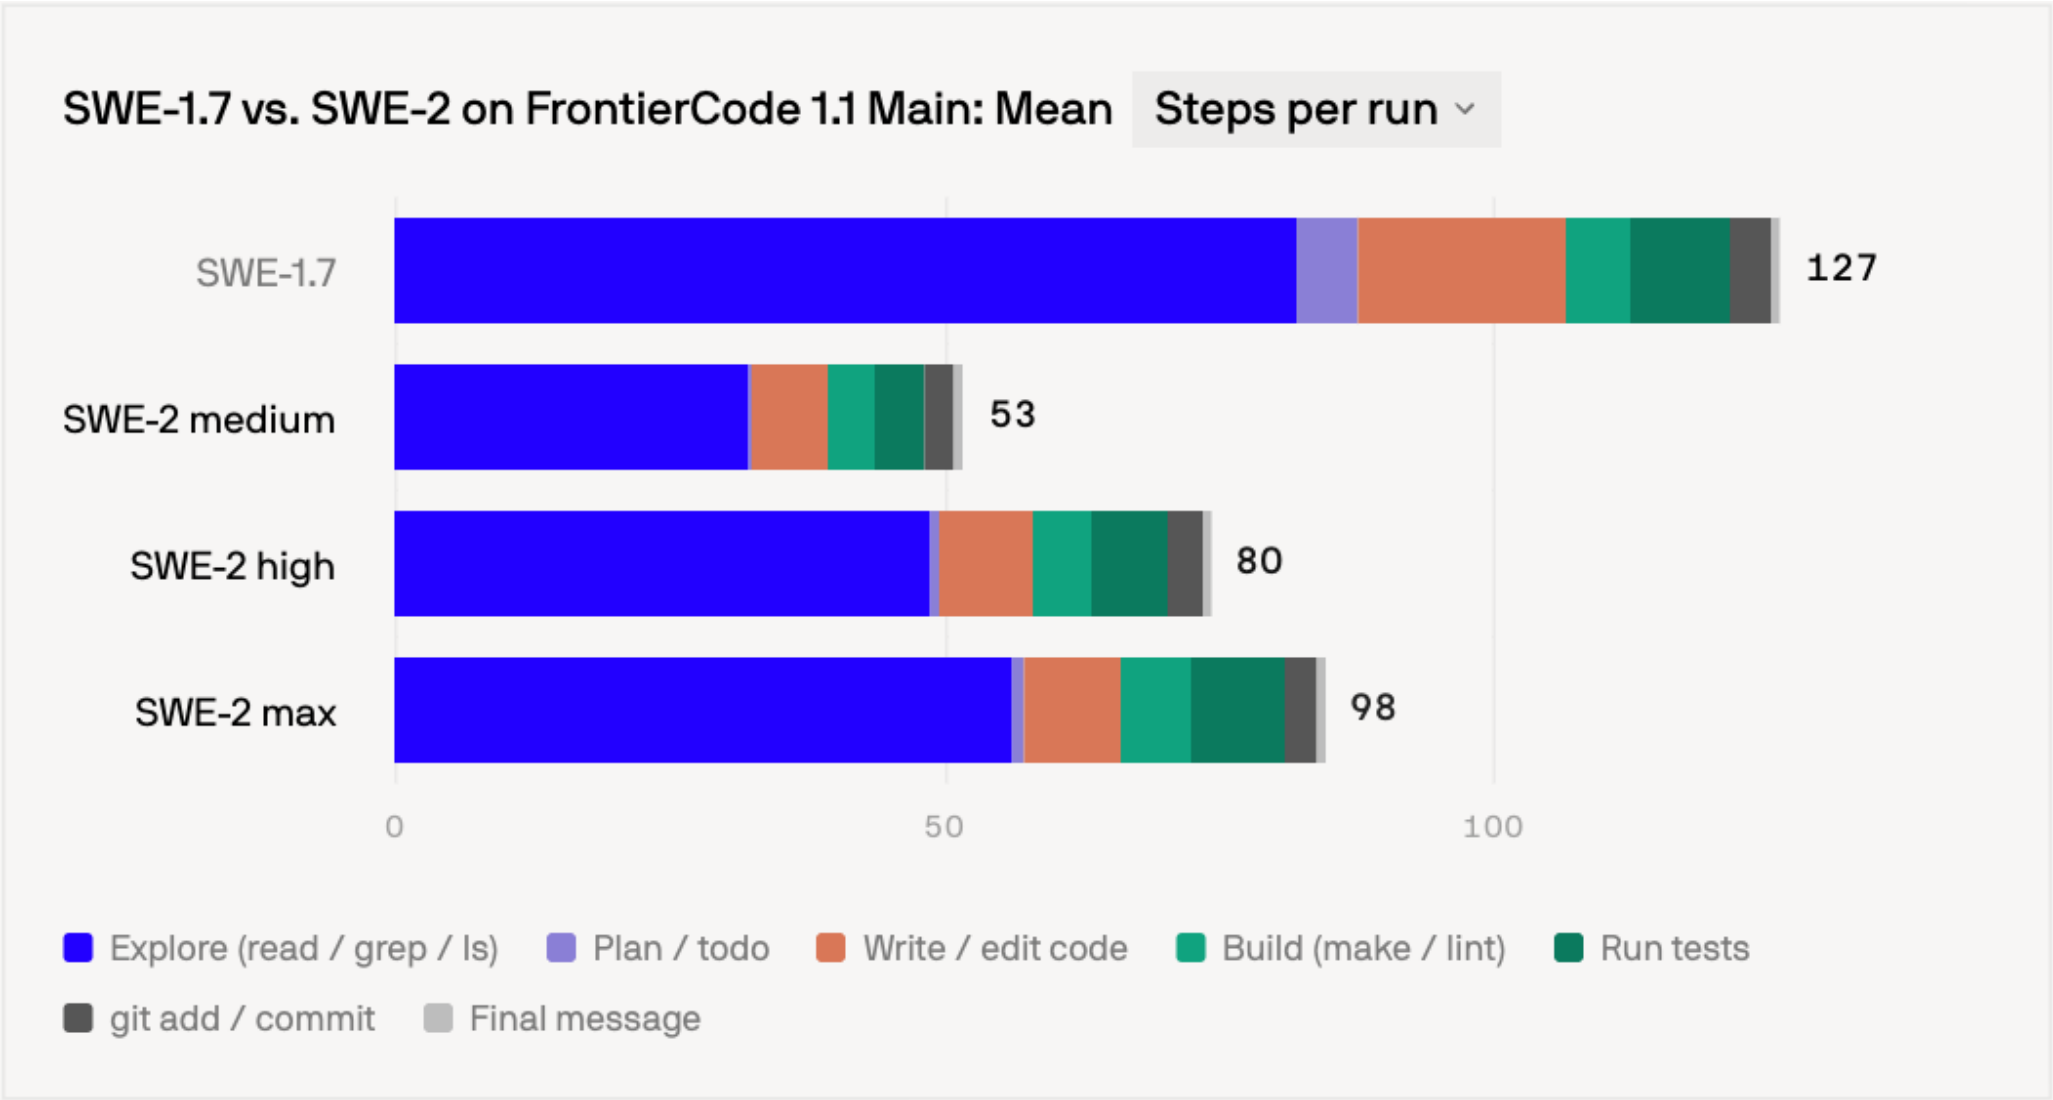
Effect of customer and language framing on vulnerability



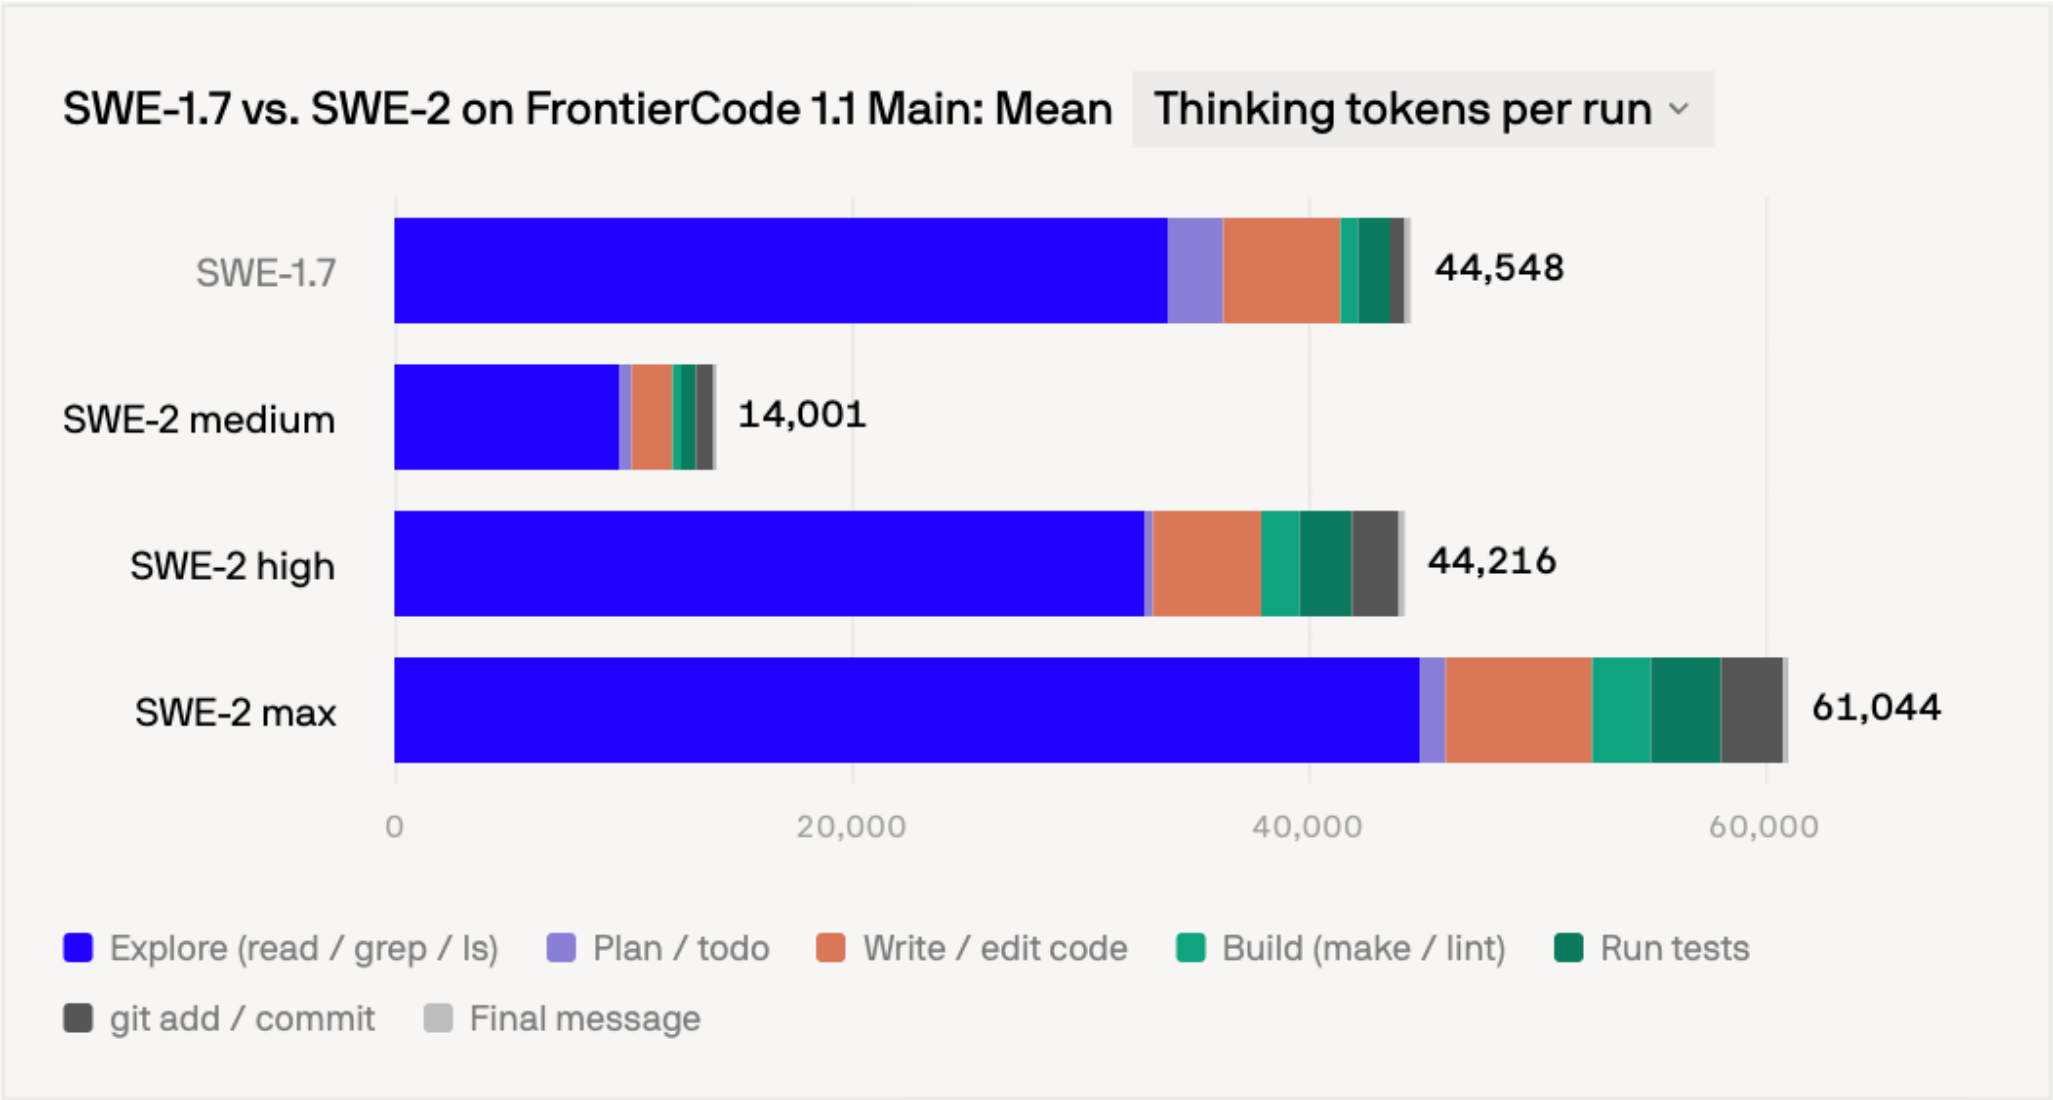
Mean judged vulnerability under each framing minus the model's overall mean; positive values are more vulnerable. Whiskers are 95% bootstrap percentile intervals over the ten tasks.



Mean metric over all 100-task FrontierCode 1.1 Main tasks, using three runs per task per model and grouped by the tools each step calls.



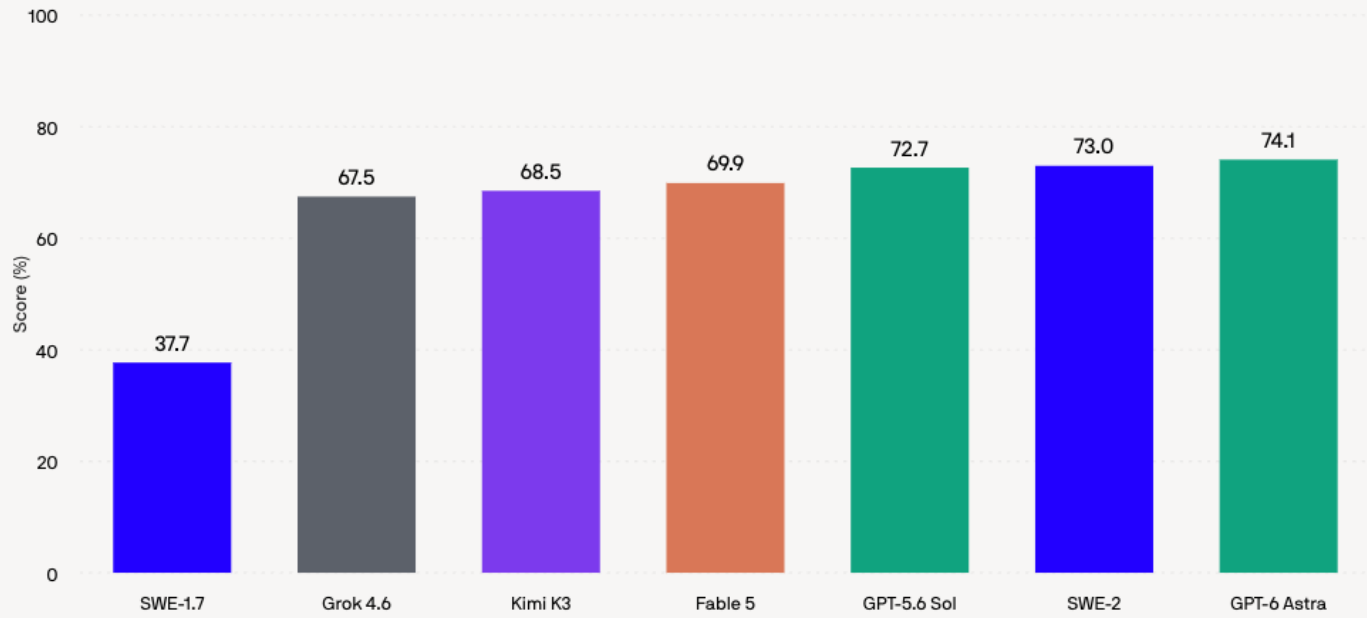
Mean metric over all 100-task FrontierCode 1.1 Main tasks, using three runs per task per model and grouped by the tools each step calls.



Mean metric over all 100-task FrontierCode 1.1 Main tasks, using three runs per task per model and grouped by the tools each step calls.

DeepSWE 1.1

Score



Score is a weighted aggregate of the rubric items. Solutions that don't pass blocking criteria receive 0.

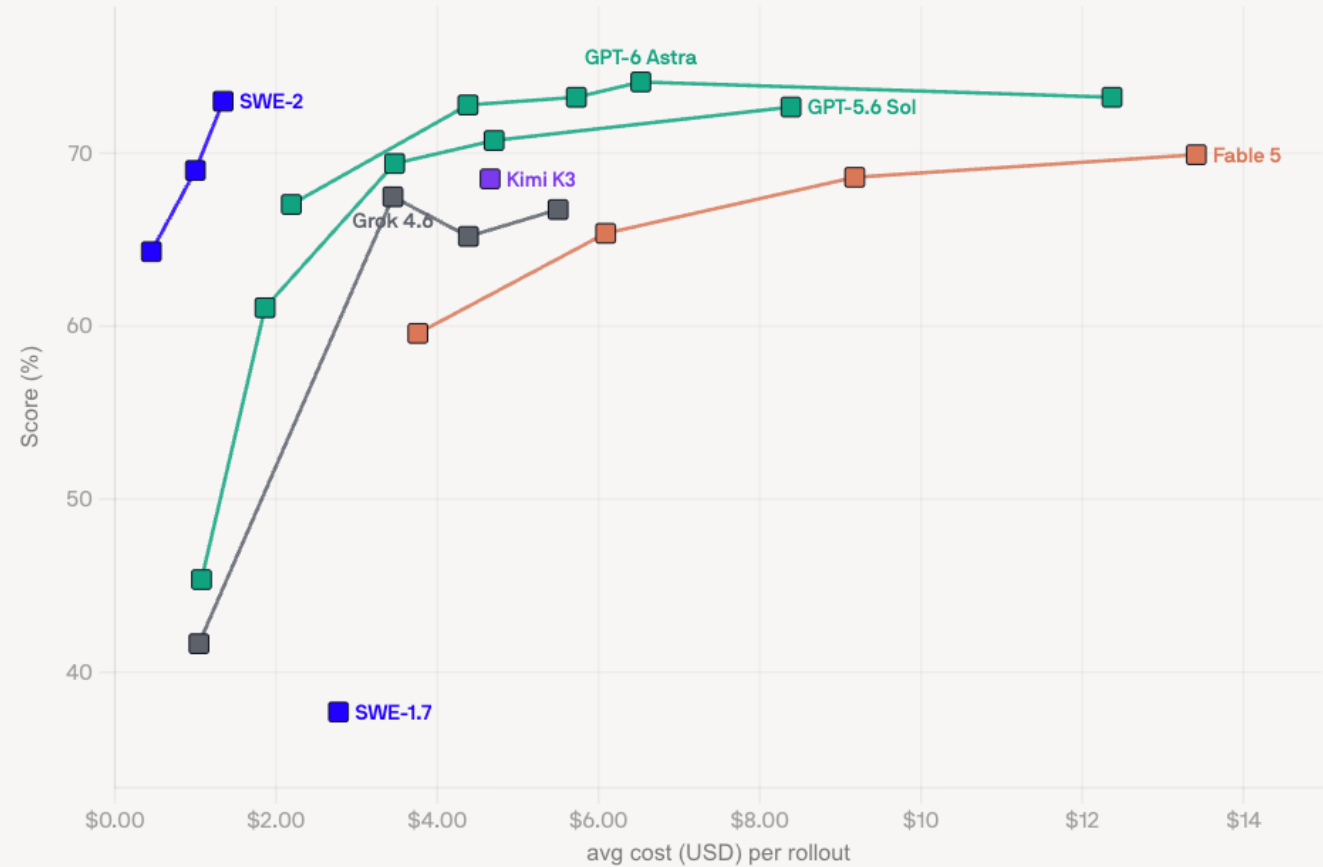
DeepSWE 1.1: the full 113-task benchmark.

DeepSWE 1.1

Score

vs

Cost (\$)



Score: the fraction of trials whose committed patch passes the task's test suite.

Cost (\$): the mean USD spend per rollout.

